

S. Towards completion

- Particularly suited for bath systems
- Extension for interactive systems:
- Jobs removed for running are put back into the back of the queue
- Starvation free** as long as earlier jobs are bounded

b. Shortest Job First

- p1: 200 transactions
- p2: 10000
- p3: 1200

- Processes are ordered by total CPU time used
- Jobs that run for less time will run first
- Reduces average waiting time if number of processes is fixed

d) Potential for starvation

3.3 Co-operative Multitasking - Round Robin with Voluntary Scheduling

Currently running processes cannot be suspended by the scheduler

Processes must volunteer to give up CPU time

3.4 Pre-emptive Multitasking - Round Robin with Timer, Shortest Remaining Time

a. Round Robin with Timer:

- Each process is given a fixed time slice ci
- After time ci , scheduler is invoked and next task is selected on a RR basis

b. Shortest Remaining Time

- Pre-emptive form of SJF
- Processes are ordered according to remaining CPU time left

Currently running processes can be force suspended by the scheduler

3.5 Real-Time Multitasking - Rate Monotonic Scheduling, Earliest Deadline First Scheduling

Processes have fixed deadlines that must be met

If don't meet the deadline:

- Hard Real Time Systems: System fails
- Soft Real Time Systems: Mostly just an inconvenience. Performance of system degraded.

4. Scheduling in Linux

Processes in Linux are dynamic: New processes can be created with `fork()`. Existing processes can exit.

Scheme ensures **no starvation of lowest priority processes.**

- How process priorities are calculated:
 - Priority = base + f(nice) + g(cpu usage estimate)
 - f() = priority adjustment from nice value
 - g() = Decay function. **Processes that have already consumed a lot of CPU time are downgraded**
- I/O boost:

Rationale:

Tasks doing read() has been waiting for a long time. May need quick response when ready

Blocked/waiting processes have not run much

Applies also to interactive processes - blocked on keyboard/mouse input

Lecture 14 Interprocess Communications

- Race Conditions:** Race condition occur when two or more processes attempt to access shared resources: Global variables, memory locations, hardware registers, CPU time

- Critical Section:** When a running process is reading/updating global variable (which can lead to race condition), it is within its **"critical section"**
- The Producer/Consumer Problem:** Consumer gets preempted, eventually no one is awake, deadlock
- Semaphores**
A semaphore is a special lock variable that counts the number of wake-ups saved for future use
A value of '0' indicates that no wake-ups have been saved
Two atomic operations on semaphore
a. **DOWN, TAKE, PEND or P:**
If the semaphore has a value > 0, it is decremented and the DOWN operation returns. **If the semaphore is 0, the DOWN operation blocks**
b. **UP, POST, GIVE or V**
If there are any processes blocking on a DOWN, one is selected and waken up. Otherwise UP increments the semaphore and returns.

4.1 Mutual Exclusion using Semaphore

When a **semaphore's counting ability is not needed**, we can use a simplified version called "mutex"

1 = Unlocked, 0 = Locked

Two processes can then attempt to DOWN the semaphore

Only one will succeed. The other will block

When the successful process exits the critical section, it does an UP to wake up others

4.2 Using Semaphores in Producer/Consumer Problem:

For two n-step (in assembly) process race conditions, the total possible cases are $n(n+1)/2$

```
#define N 100
typedef int semaphore;
semaphore mutex = 1;
semaphore empty = N;
semaphore full = 0;
```

4.3 Deadlocks with Semaphores

Our producer/consumer solution swapped the semaphores for empty/full with the mutex semaphore, the potential deadlock occurs

5. Monitors and Conditional Variables

Monitors achieve mutual exclusion, but we also need other mechanisms for coordination. E.g. in our producer/consumer problem, **mutual exclusion is not enough to prevent the producer from proceeding when the buffer is full.**

We introduce "condition variable". One process WAITs on a condition variable and blocks, until... Another process SIGNALs on the same condition variable, unblocking the WAITing process

6. Barriers

A "barrier" is a special form of synchronisation mechanism that works with groups of processes rather than single processes. The idea of a barrier is that all processes must reach the barrier (signifying the end of one phase of computation) before any of them are allowed to proceed. **Can be implemented by semaphores.**

Lecture 15 Memory Management

1. Logical and Physical Addresses

Logical addresses: These are addresses as "seen" by executing processes code

Physical addresses: These are addresses that are **actually sent to memory to retrieve data or instructions**

2. Base and Limit Registers

Base Register:

This contains the starting address for the program
All program address are computed relative to this register
Limit Register:

This contains the length of the memory segment
These registers solve both problems:

- We can **resolve address conflicts by setting different values in the base register**
- If a program tries to access memory **below the base register value or above the (base + limit) register value, a "segmentation fault" occurs**

3. Partitioning issues: Fragmentation

Internal fragmentation:

Partition is **much larger than is needed**
Cannot be used by other processes

Extra space is wasted

External fragmentation:
Free memory is broken into small chunks by allocated memory

Sufficient free memory in TOTAL, but individual chunks insufficient to fulfil requests

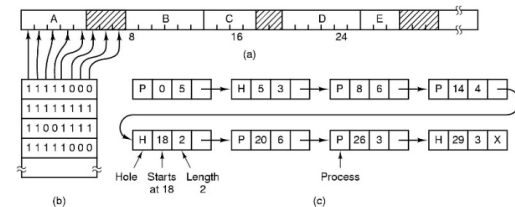
4. Managing Free Memory

Two approaches:

- Bit maps
- Free/Allocated List

In either approach, **memory is divided up into fixed sized chunks called "allocation units"**

Common sizes range from several bytes (e.g. 16 bytes) to several kilo-bytes
Each "tick mark" in figure (a) represents the boundary of an allocation unit



4.1 bitmaps

Each bit corresponds to an allocation unit

0 = free, 1 = allocated

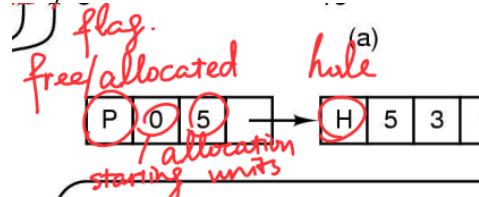
Calculation questions:

① **Max amount of memory manageable = len(bitmap) in bits × n bytes/bit**
② **IOU size 1 bit ~ n bytes**

malloc: allocate in bytes
When allocating in bytes, **calculate their corresponding unit(s) in AU** (round to ceiling) and reflect in bitmap.

③ **Internal Fragmentation = (n of IOU × size of IOU) - actual num of bytes**

4.2 Free/Allocated List



5. Allocation Policies

5.1 First-Fit

- Scan through the list/bit map and find the first block of free units that can fit the requested size
- Fast, easy to implement

5.2 Best Fit

- Scan through the list/bit map to find the smallest block of free units that can fit the requested size
- Minimise waste
- It can lead to scattered bits of tiny useless holes

5.3 Worst Fit

- Find the largest block of free memory
- Can reduce the number of tiny useless holes

5.4 Buddy Allocation

Binary splitting
Half of the block is allocated
The two halves are called "buddy blocks"
Can combine again when two buddy blocks are free

Lecture 17 File Systems

1. I-node in UNIX

Owner id
File type
Protection information
Mapping to physical disk
Time of creation, last use,
Reference counter

2. Linked List Allocation -

Linked list implemented 4 types:

- Block number (part)
- EOF byte (end of line)
- FREE byte
- BAD byte

Some numbers are block block numbers for EOF, F
Blocks for file linked together
3->5->8

Free blocks indicated by FREE entry

Bad blocks (unusable blocks due to disk error) marked in FAT entry: BAD cluster value

2.1 FAT16

Total # of blocks = 2^{16} (- some reserved blocks)

Max size = Total # of blks * size (bytes/block)

FAT32

FAT32

■ Increase FAT size to 28 bits, cluster numbers 28-bits

Inode calculation questions:
1. **pointer size = 8 bit / bytes**
 $n = \text{Max \# of pointers} = \frac{\text{pointer size}}{\text{block size}}$
2. **Max. size of a single file =**
 $\sum_{i=1}^n \text{ndirect} \times l + \sum_{i=1}^n n_i \times n_p \times l$
where block size = 2 bytes
E.g. $10 \times 2048 + 512 \times 2048 + 512 \times 2048$
direct 1 bit 1 bit 1 bit
nd: 10 bit-1 bit-1
3. **Max partition possible**
 $2^{\text{len(bitmap) in bits}} \times l$